

# Distributed Probabilistic Counting on ClickStream data

Michael Vogiatzis  
Software Engineer  
@mvogiatzis

# About me

- Software Engineer in London
- Interested in
  - All things scaling
  - Real-time processing
  - Open Source
  - Good whiskey
- Msc in CS, University of Edinburgh

# What is ClickStream data?

- Recording of user clicks on a web page or software app
- User actions are logged for further analysis
- USA Gov gives (some of) them for free!

# ClickStream use cases

- User behavior analysis
  - Join with CRM, identify users, target ads
- Market research
- Software testing
- UX improvement
- Fraud detection

# Problem

Find the most visited pages and/or most active users in a set of websites



# Problem

Find the most visited pages and/or most active users in a set of websites

count++

# Problem Solved !




@mvogiatzis

# Solved?

Now count clicks on the **Amazon** website in **real-time** including **spikes**.





@mvogiatzis

# Randomised Algorithms

## Motivation

What if storage / processing times grows exponentially?

# Randomised Algorithms

## Motivation

What if storage / processing times grows exponentially?

- Computational & space efficiency

# Randomised Algorithms

## Motivation

What if storage / processing times grows exponentially?

- Computational & space efficiency
- Scaling



# Randomised Algorithms

## Motivation

What if storage / processing times grows exponentially?

- Computational & space efficiency
- Scaling
- No need for *exact* counts
  - Perhaps top-N most frequent ones



# Randomised Algorithms

## Motivation

- Randomised approaches are often the most efficient approach to many classes of demanding problems
- Trade-off between error rate and performance
- Bloom filters, Locality sensitive hashing, **probabilistic counting**, etc.

# Probabilistic Counting

- Often referred to as “Approximate counting”
- “The approximate counting algorithm allows the counting of a large number of events using a small amount of memory.”
  - Wikipedia
- A randomised counting approach

# How It Works

- Every time a new event comes update the counter with probability  $2^{-f}$ 
  - where  $f$  = current\_frequency
- Meaning: Update the counter only if a random generated number (sampled uniformly between 0 and 1) is less than  $2^{-f}$

# How It Works

- Every time a new event comes update the counter with probability  $2^{-f}$ 
  - where  $f$  = current\_frequency
- Meaning: Update the counter only if a random generated number (sampled uniformly between 0 and 1) is less than  $2^{-f}$
- We now need  $\log(\log(f))$  bits per counter instead of  $\log(f)$ 

This counts in log-space.



# Example

If  $\text{Random\_Num} < 2^{-\text{prev\_count}}$   
Count++;

Stream	Prev Count	Random Num	Curr Count
"dummy"	0	0.6	1
"dummy"	1	0.7	1
"dummy"	1	0.3	2
"dummy"	2	0.2	3
"dummy"	3	0.5	3



# Exact results

- We may end up over-counting or under-counting
- $E[2^f - 1]$  gives the estimated real count
- Decrease the error rate by using a smaller base than 2
  - Updates more frequently

# Scale with Storm

- Single machine counting may not be enough

# Scale with Storm

- Single machine counting may not be enough
- Use multiple machines for *load balancing* and *fault-tolerance*

# Storm


- Distributed real-time computation system
- Fault-tolerant
- Fast
- Scalable
- Guaranteed message processing
- Open source
- Multi-language capabilities

# Elements

- Streams
  - Set of tuples
  - Unbounded sequence of data
- Spout
  - Source of streams
- Bolt
  - Application logic
  - Streaming aggregations, joins, DB ops



# Topology




# Data source

- 1.USA.gov URLs are created whenever anyone shortens a .gov or .mil URL using bitly.
- <http://developer.usa.gov/1usagov>

# Data format

```
{  
  "a": USER_AGENT,  
  "c": COUNTRY_CODE, # 2-character iso code  
  "nk": KNOWN_USER, # 1 or 0. 0=first time we've seen this browser  
  "g": GLOBAL_BITLY_HASH,  
  "h": ENCODING_USER_BITLY_HASH,  
  "l": ENCODING_USER_LOGIN,  
  "hh": SHORT_URL_CNAME,  
  "r": REFERRING_URL,  
  "u": LONG_URL,  
  "t": TIMESTAMP,  
  "gr": GEO_REGION,  
  "ll": [LATITUDE, LONGITUDE],  
  "cy": GEO_CITY_NAME,  
  "tz": TIMEZONE  
  "hc": TIMESTAMP OF TIME HASH WAS CREATED,  
  "al": ACCEPT_LANGUAGE  
}
```

it's DEMOtime!



# The End

- Code on GitHub
  - [.github.com/mvogliatzis/probabilistic-counting](https://github.com/mvogliatzis/probabilistic-counting)
- Follow me: [@mvogliatzis](https://twitter.com/mvogliatzis)
- E-mail me: [michael@micvog.com](mailto:michael@micvog.com)
- Read my blog: <http://micvog.com/>



# Q & A

